

OSIPI · GSOC 2026 · ASL QUALITY CONTROL

# Kidney QC the short version

---

Extending the ASL QC ToolBox from brain to kidney

Agnik Misra

KIDNEY\_QC\_ONE\_PAGER.md · 1,339 words

reconciled against osipy-qc · 20 registered checks · 253 tests

## In one paragraph

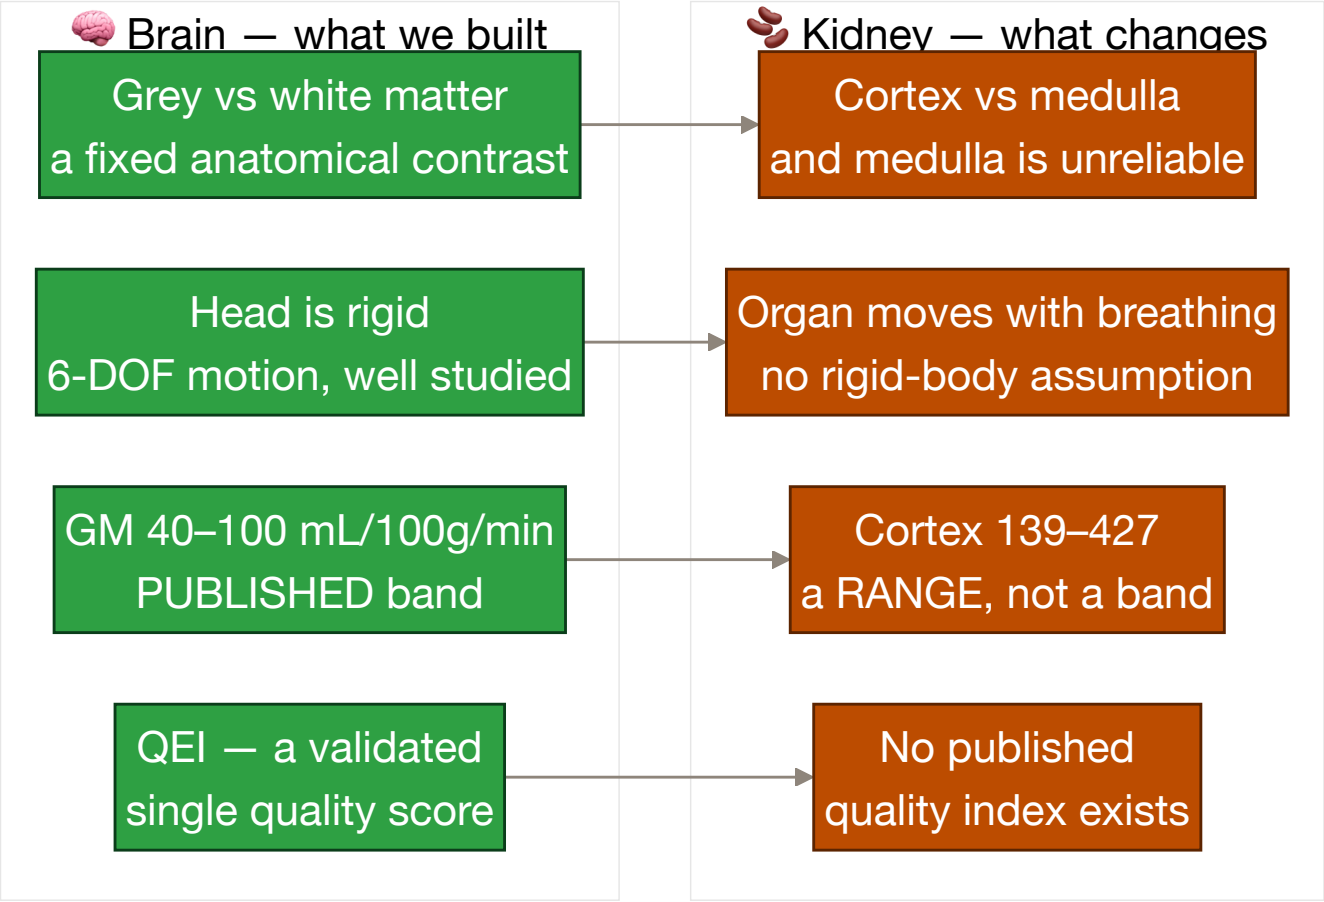
The brain toolbox grades a CBF map against grey matter. The kidney has no grey matter. It has a **cortex and a medulla**, it **moves with every breath**, and it is perfused roughly **five times harder** than the brain. Most of the maths survives that move; almost none of the *thresholds* do. The single most important finding from reading the literature is a negative one, and it shapes everything below.

### The finding that shapes the design

**Nery et al. 2020** (MAGMA, PARENCHIMA COST Action) is the renal equivalent of the ASL White Paper — 59 consensus statements, and I read every one. It contains **zero numeric quality thresholds**. No tSNR cutoff, no CoV cutoff, no motion limit, no QEI equivalent.

So the published material is all about *how to acquire* the scan, and none of it is about *whether the map is good*. v1 grades the map and leaves protocol conformance out of scope, which means every threshold we ship is labelled **UNCALIBRATED**. That is a publishable statement about the field, not a gap in our work.

🧠 → 🦋 Why the kidney is harder



**The perfusion numbers do not converge.** Published healthy cortical perfusion spans **139–427 mL/100 g/min**. The dominant driver is not biology — it is the **labelling scheme**: **FAIR reads 1.5–1.8× higher than pCASL in the same subjects** ( $362 \pm 57$  vs  $201 \pm 72$ ). Field strength adds ~10%, age another ~15–20%. A band that is not conditioned on labelling scheme and field strength is measuring the protocol, not the patient.

⚖️ **Two consensus rules we get for free**

These are the only *published mandates* that map directly onto our report, and both are cheap to honour:


| RULE                                                                                                       | AGREEMENT   | WHAT IT MEANS FOR US                                               |
|------------------------------------------------------------------------------------------------------------|-------------|--------------------------------------------------------------------|
| <b>R10.1</b> — report <b>cortical</b> RBF (not whole-kidney), <b>separately for left and right</b>         | <b>100%</b> | our report schema is decided for us: per-kidney, cortex-restricted |
| <b>R10.2</b> — medullary values " <b>are not considered reliable with current measurement approaches</b> " | <b>89%</b>  | 🚫 the cortex:medulla ratio <b>cannot be a perfusion verdict</b>    |

R10.2 is the one that hurts. The cortex:medulla ratio looked like the obvious kidney analogue of the brain's GM/WM ratio — scale-free, contrast-based, exactly the kind of check that survives a calibration error. Consensus says the medullary half of it is not trustworthy. So it ships as a **segmentation-integrity flag** (INFO/WARN, never FAIL): if the ratio inverts, the *masks* are probably swapped — which is a real and useful thing to catch, just not the thing we hoped for.

---



## The proposed check set



**Protocol conformance is deliberately out of scope for v1.** Nery's 59 statements are the only published, directly checkable requirements in renal ASL — labelling scheme, timings, geometry, readout, quantification constants — and v1 does not grade any of them. That is a real cost, stated plainly: it leaves the design resting almost entirely on **uncalibrated** thresholds. The boundary is the same one the brain tool draws — grade the data, not the acquisition parameters.

**K1 is deliberately empty** for a different reason: there is no renal QEI, and inventing one would be the single least defensible thing we could do.

**K7 carries the one genuinely implementable published rule family** — subtraction-outlier rejection. Respiratory motion is the dominant renal artefact, and unlike the brain there is no rigid-body assumption to lean on.

## What you must bring

The kidney has no equivalent of "just run BET on it" — **a perfusion map alone buys almost nothing**, because every meaningful metric is ROI-restricted and the ROIs must come from outside.

| TIER      | SUPPLY                                                 | UNLOCKS                                                                              |
|-----------|--------------------------------------------------------|--------------------------------------------------------------------------------------|
| <b>B0</b> | RBF map only                                           | almost nothing — histogram shape, implausible-value fraction                         |
| <b>B1</b> | + whole-kidney masks, <b>left and right separately</b> | negative fraction, left-vs-right consistency, mask integrity                         |
| <b>B2</b> | + <b>cortex mask per kidney</b>                        | <b>cortical RBF — the quantity R10.1 mandates</b>                                    |
| <b>B3</b> | + medulla mask (or a T1 map)                           | cortex:medulla ratio, as an integrity flag only                                      |
| <b>A0</b> | raw 4D + kidney masks                                  | outlier rate, respiratory displacement, control/label order                          |
| <b>A1</b> | + M0                                                   | M0 checks, perfusion signal as % of M0                                               |
| <b>A2</b> | + acquisition metadata                                 | four existing checks sharpen; <b>no new check</b> — v1 grades no protocol parameters |
| <b>A3</b> | + respiratory trace                                    | gating efficiency, motion-vs-outlier correlation                                     |

## What transfers from the brain code

| BRAIN CHECK                                                  | VERDICT                                                              |
|--------------------------------------------------------------|----------------------------------------------------------------------|
| Negative-voxel fraction, histogram                           | ✅ <b>transfers as-is</b> — non-physical is non-physical in any organ |
| Schema, data type, M0 presence / TR / background suppression | ✅ <b>transfers</b> , with renal metadata fields                      |

| BRAIN CHECK                       | VERDICT                                                                                                                            |
|-----------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| CBF level, left-right consistency | ◆ <b>new bands</b> — and they must be conditioned on labelling scheme                                                              |
| GM/WM ratio → cortex:medulla      | ◆ <b>reworked</b> — demoted to an integrity flag by R10.2                                                                          |
| Motion (FWD/DVARS)                | ◆ <b>reworked</b> — the organ moves, the head does not                                                                             |
| <b>QEI</b>                        | ✗ <b>does not transfer</b> — it is built on GM/WM/CSF probability maps and a grey-matter spatial prior. No renal equivalent exists |

**The architecture change this forces:** `cfg.organ` already exists in the config but **no check reads it** — it is inert today. Kidney is the reason to make it real, so checks, thresholds and profiles become organ-aware rather than kidney being bolted on beside brain.

## ? For the mentors and the renal specialist

1. **R10.2 says medullary RBF is unreliable — do you agree the cortex:medulla ratio should be an integrity flag rather than a quality verdict?** It is the most useful-looking metric we have and consensus undercuts it.
2. **Should cortical bands be conditioned on labelling scheme?** FAIR vs pCASL differ 1.5–1.8× in the same subjects. One band cannot serve both, and pretending otherwise grades the protocol.
3. **Is there an unpublished or in-progress renal quality index we should build toward** rather than leaving K1 empty?
4. **What is realistically available in practice** — do sites have cortex and medulla masks, or is B1 (whole-kidney only) the honest default tier?

**Full detail:** `BRAIN_TO_KIDNEY.md` → `KIDNEY_AS_L_EXPLAINED.md` → `KIDNEY_QC_DESIGN.md` **Read first:** `papers/ner2020_renal_asl_consensus.pdf` 51 papers · 68 highlighted source pages